

RAG Blog Series

Code companion for the RAG Blog Series. Each folder corresponds to a blog post with its own README, requirements, and code, guiding you from keyword search to scalable, production-ready RAG systems.

RAG Blog Series — Code Companion

Code companion for the **RAG Blog Series**.

Each folder corresponds to a blog post with its own README, requirements, and runnable code.

The series walks you through the journey of **Retrieval-Augmented Generation (RAG)** — from simple keyword search all the way to scalable, production-ready RAG systems.

Blog Series Index

- [Blog 1 – RAG Demystified: A Beginner’s Guide to What It Is and Why It Matters](#)

Learn why LLMs hallucinate, why retrieval matters, and build a toy keyword-based Q&A system.

- [Blog 2 – Speaking in Vectors: A Beginner’s Guide to Embeddings and Semantic Search](#)

Understanding embeddings, cosine similarity, and FAISS-the building blocks of smarter RAG systems.

(More coming soon: evaluation, monitoring, observability, deployment...)

How to Use This Repo

Each blog folder is independent.

To run code for a specific blog:

```
```bash cd 1_rag_demystified
python -m venv .venv
```

### On Linux/Mac:

```
source .venv/bin/activate
```

# On Windows:

`.venv pip install -r requirements.txt python toy_rag.py # or the relevant script for that blog`